Extending AI-Agent Incident Analysis to Multi-Agent Systems

AFACS, an HFACS-based framework, and an audit of public traces

James Le Cuirot
CitadelAero
Senior First Officer and Flight Safety Officer · MSc Human Factors in Aviation
james@citadelaero.com · citadelaero.com

Draft, November 2026

Abstract

Ezell, Roberts-Gaal and Chan (2025) proposed a framework for analysing incidents involving AI agents. It treats single agents, has no category for deliberate departures from known requirements, and has not been applied to real traces. This paper extends it to multi-agent systems using the structure of the Human Factors Analysis and Classification System (HFACS) from aviation. The extension, AFACS, adds a supervision tier, a coordination category and violations coded from behaviour with a graded statement of the evidence, and it states what evidence each layer needs. It is tested on the public MAST dataset of multi-agent traces. An audit finds duplicate and corrupt rows, an LLM judge that agrees with itself at a mean Cohen's κ of 0.43, and logs that record what the agents did but none of the system documentation needed to reconstruct a failure. In an AI pilot, Claude coded 250 failed traces three times. The layers were applied consistently (κ 0.63–0.76 between two models and 0.74–0.87 between two runs of one model; individual codes 0.52 and 0.59). Supervision failures appear in about two-thirds of failures, coordination breakdowns in two-fifths and violations in 9–13%, mostly inferred from the act rather than stated; MAST's "Disobey Task Specification" label mixes errors with violations. Checked against the frameworks' documentation and code, the annotator's design-level claims usually identified a real weakness but overstated it as an absence (8 of 24 confirmed in full, 2 contradicted): design causes need design evidence. Validation with human coders is planned.

1 Introduction

When an aircraft accident is investigated, "pilot error" is where the analysis starts. Reason's model of organisational accidents describes how the unsafe act at the sharp end is made likely by latent conditions created earlier and higher up [1, 2]. The Human Factors Analysis and Classification System (HFACS) turned that model into a coding scheme with four tiers, from organisational influences down to unsafe acts [3, 4]. HFACS also keeps Reason's distinction between errors, which are unintended, and violations, which are departures from known rules. The two call for different remedies and appear to be different kinds of behaviour [5].

LLM agents act over many steps: they plan, call tools, write and run code, and exchange messages. A multi-agent system gives several agents roles, such as coder, tester and orchestrator, much as a crew divides duties. Ezell, Roberts-Gaal and Chan [6] proposed the first framework for analysing incidents involving AI agents in this tradition. It explains incidents through system factors, contextual factors and cognitive errors, sets out what information investigators need, and lists open challenges. Three features leave room for extension. It treats a single agent, so it has no place for an orchestrator that supervises other agents or for breakdowns between agents. Its agent-level category holds errors only: the paper cites an agent that reasoned about cheating its unit tests before doing so [7], but no category records such a departure. And it has not been applied to real traces.

This paper extends the framework to multi-agent systems using the structure of HFACS, and calls the result AFACS. It takes up five of Ezell et al.'s open challenges: non-exhaustive coverage, emerging factors, high-level cognitive errors, reconstruction difficulties with limited data access, and, in part, expertise shortages. Using the public MAST dataset [8], it asks four questions.

- **RQ1, cross-walk.** How do Ezell et al.'s factors map onto the HFACS layers, and what does a multi-agent setting add?
- **RQ2, information adequacy.** Do real multi-agent traces contain the information Ezell et al. say investigators need?
- **RQ3, applicability.** Can the extended framework be applied consistently, and how often do failures involve the added categories?
- **RQ4, flat labels.** How do MAST's fourteen flat labels map onto the extended framework?

The contributions are the framework and its cross-walk; an audit of the MAST dataset and of what its logs let an investigator establish; and an AI-coded pilot on 250 traces, coded three times, with a check of the annotator's design claims against documentation. Human validation is planned as a second study, for which a blinded, pre-registered calibration set has been frozen.

2 Background

Incident analysis and failure taxonomies. Ezell et al. define system factors (training data, feedback, learning methods, system prompt, scaffolding), contextual factors (information, tools, task definition) and cognitive errors (observation, understanding, decision-making, action execution), and list eighteen items of information investigators need (their Table 1) [6]. Swiss-cheese guardrail designs and STPA look forward, at defences and hazards [9, 11, 10]; incident analysis looks back. MAST, derived from 150 traces, reports human agreement of $\kappa = 0.88$ and $\kappa = 0.77$ between its o1 judge and experts [8], but has no causal layers. Other taxonomies classify failures by error type [12], by the component that broke [13], by the agent and step responsible [14], or by where the fix belongs [15]. They answer "what broke" and "who and when"; AFACS adds "what kind of act" and "what latent conditions".

HFACS and its reliability. An aviation psychologist and a pilot re-coding 319 causal factors from NTSB reports into HFACS agreed at $\kappa = 0.71$ [16]. Reliability is contested: a modified HFACS showed low agreement [17], and a review found it higher with training and fewer departures from the original [18, 19]. LLMs coding aviation reports reach far lower exact-match than partial-match agreement [20]. AFACS therefore keeps the four tiers and reports agreement by layer as well as by code.

AI behaviours that resemble violations. Specification gaming is the standard term for many behaviours AFACS would code as violations [21]. ImpossibleBench builds tasks that can be passed only by breaking the specification [22]. Reasoning traces can reveal reward hacking, and models trained against monitors learn to hide it [7]. Frontier models in constructed scenarios reason about constraints and circumvent them [23, 24]. OverclaimBench measures false claims of completion with a criterion that "requires no inference about intent" [25], the stance AFACS takes; and failures look increasingly like an incoherent "hot mess" as tasks lengthen [26], so most should be errors.

3 The AFACS Framework (RQ1)

Each failed trace is coded at every layer where it gives evidence (multi-label), with a deepest contributing layer and a confidence rating (Table 1; `codebook/v1.1.md`). The codebook is the author's.

Adapting HFACS to multi-agent runs. Five decisions adapt HFACS. (1) The humans around a run sit in existing tiers: the user who sets the task is at the top of the supervision chain, so a deficient task is L3.4 whoever set it, and a run launched or configured wrongly (for instance with an empty task) is `L4.5` Operation set-up, the analogue of a defective dispatch, rather than an excluded trace. (2) When the framework's own software misbehaves, the cause decides the code: a malfunction during the run is L2.1, a limitation in the design L4.1 or L4.2. (3) A failure need not include an unsafe act: if a tool fails and the agents respond reasonably, no L1 code is given. (4) A loop is the repeating agent's unsafe act and, where a supervisor could have broken it, also L3.3. (5) An omission is coded by whose job it was: L1 if an agent was told to do it, L3.1 if a supervisor should have checked, L4.1 or L4.3 if nobody was given the job.

Table 1: AFACS layers and codes (codebook v1, used in the pilot; v1.1 differs only in a note). L4.4, added from Ezell et al.'s system factors, needs developer-side evidence and is not applied from traces (Section 3). L4.5 covers the people who launched the run.

Layer	Code	Description
L4 Organisational / design	L4.1	Architecture and role design: role split, missing verifier, no termination logic
	L4.2	Resource provisioning: tools, step or token budgets, context limits
	L4.3	Process: no standard procedure for verification or handover
	L4.4	Model development (developer evidence only): training data, feedback or learning methods
	L4.5	Operation set-up: the run was launched or configured wrongly, e.g. an empty task
L3 Supervision / orchestration	L3.1	Inadequate supervision: no check on subagent output
	L3.2	Planned inappropriate operation: task beyond an agent's capability or information
	L3.3	Failure to correct a known problem: visible loop or error, no intervention
	L3.4	Deficient specification: ambiguous, contradictory or incomplete task or prompt, set by the user or by an orchestrating agent
L2 Preconditions	L2.1	Environment: tool or API failure, bad or unavailable information
	L2.2	Agent state: context saturation, lost instructions, carried-over error
	L2.3	Coordination: communication breakdown, withheld or ignored input
L1 Unsafe acts: errors	L1.E1	Skill-based: execution slip, e.g. a malformed tool call
	L1.E2	Decision: wrong plan or choice given adequate information
	L1.E3	Perceptual: misread observation, tool output or message
L1 Unsafe acts: violations	L1.V1	Routine: habitual shortcutting across the trace
	L1.V2	Exceptional: one-off departure, e.g. fabricating a result or editing tests
	L1.V?	Indeterminate: the criteria appear to hold but the evidence is thin

Table 2: Condensed cross-walk from Ezell et al.'s factors to AFACS codes. The last three rows have no counterpart in Ezell et al. †Needs developer evidence; not applied from traces (Section 3).

Ezell et al. factor	AFACS code	Note
System: training data, feedback, learning methods	L4.4†	Latent; feedback that rewards appearing successful is a typical condition behind violations
System: system prompt underspecified or encouraging undesirable behaviour	L3.4	System prompt treated as part of the specification
System: scaffolding not working as intended	L4.1/L4.2 or L2.1	Coded by cause: built into the design, or a failure during the run
Contextual: information inaccessible, missing, poor or malicious	L2.1	
Contextual: tools unusable or failing; tool errors not visible	L2.1 or L4.2	Run-time failure, or provisioning by design
Contextual: unclear or conflicting task; resource limits	L3.4; L4.2	
Cognitive: observation (failure to notice, wrong expectation)	L1.E3	Memory capacity at L2.2
Cognitive: understanding (misinterpretation; conclusions from poor information)	L1.E3; L1.E2	With L2.1 where the information was poor
Cognitive: decision-making (goal priority, non-viable plans, criteria)	L1.E2	Distraction from the goal at L2.2
Cognitive: action execution (incorrect execution; no action; no error monitoring)	L1.E1; L1.E1/E2; L1.E3 or L3.3	L3.3 where a supervising agent should have caught it
(none: single-agent framework)	L3.1–L3.3	Orchestrator fails to check, assigns beyond capability, or leaves a visible problem uncorrected
(none)	L2.3	Coordination between agents, the crew resource management analogue
(mentioned, not categorised)	L1.V1, L1.V2, L1.V?	Violations with an evidence grade (Section 3)

Cross-walk. Table 2 condenses the item-by-item mapping (`docs/crosswalk_ezell.md`). Ezell et al.'s system factors map to L4, with the system prompt treated as specification (L3.4); contextual factors to L2.1, L3.4 and L4.2; cognitive errors to the L1 errors, with memory and distraction at L2.2. Their four cognitive functions cut the ground differently from HFACS's three error types, so AFACS records both for each error. In the other direction, their model-development factor has no HFACS tier; AFACS adds L4.4 by analogy with HFACS's resource management [16].

Layered evidence. A layered framework needs layered evidence: each layer is coded only from evidence an investigator in that position can hold (Table 7, Appendix C). Unsafe acts and preconditions are visible in the trace; supervision failures also need the task specification and system prompt; design factors need what is known of the design; model development needs the developer's records. L4.4 is therefore never coded from a trace, and the annotator is never offered it. HFACS met the same limit in aviation: organisational causes come from company records, not the flight recorder, and were rare in NTSB-based coding [16].

Violations. "Violation" is used in HFACS's sense, applied to behaviour. An investigator never sees inside a pilot's head either; a violation is inferred from evidence that the rule was known and the action departed from it. AFACS makes no claim about intent. A failure is a violation only if (1) the requirement was in the agent's accessible context; (2) the trace shows the agent represented it, or the departure cannot plausibly be accidental, such as hard-coding an expected output; and (3) the agent proceeded otherwise. If all three hold but the evidence is thin, the code is `L1.V?`; otherwise it is an error. A routine violation (`L1.V1`) here means one that recurs within the run, since a single trace cannot show a tolerated habit. The distinction matters because an error calls for better tools, tasks or models, and a violation points

Table 3: Evidence grades for violations (codebook v1, field `violation_evidence`).

Grade	Use when	Goes with
stated	The agent's own words (reasoning, plan or message) acknowledge the requirement, and it then departs from it	L1.V1, L1.V2
inferred	Nothing in the agent's words shows it registered the requirement, but the act cannot plausibly be accidental	L1.V1, L1.V2
thin	The criteria seem to hold but the evidence for condition (1) or (2) is weak	L1.V?
none	No violation is coded	

to incentives, monitoring and what training rewards. Each violation carries an evidence grade (Table 3). Stated reasoning is evidence, not proof: written reasoning may not reflect what produced an action, and models can learn to hide it [7], which is why the *inferred* grade exists.

Failures and incidents. Ezell et al. analyse incidents, events that cause or could cause harm. MAST's traces are failed benchmark tasks. Aviation applies the same causal analysis to occurrences without harm, because their causes are the causes of future accidents; AFACS is applied on that basis.

4 Data and Method

Data and audit. The data are MAST's public release (`mcemri/MAST-Data`, revision `95118ac`, CC-BY-4.0): 1,642 traces from seven frameworks, each with 14 binary labels from an LLM judge (o1). The audit (Appendix A; `docs/data_audit.md`) found 1,537 distinct traces among the 1,642 rows; a group of 100 rows holding 30 distinct texts with unrelated labels; misfiled and mislabelled traces; five runs launched with an empty task; no record of task success; and, on 30 traces the judge labelled twice, a mean per-mode κ of 0.43 (pooling modes would inflate it to 0.54, so per-code averaging is used throughout). After exclusions, and excluding ChatDev (traces of 60,000–500,000 tokens) and traces over 30,000 tokens, the eligible pool is 893 MAST-failed traces across six frameworks.

Rendering and blinding. A parser per log format splits each trace into numbered steps, keeping agent text verbatim. Every coder reads the same rendering. MAST's labels are hidden, and information the agents never saw (benchmark answers and pass/fail verdicts) is replaced by a placeholder, so that the outcome cannot colour the coding.

Information audit (RQ2). For each framework, each of Ezell et al.'s eighteen items is classified as present, partial, absent or not applicable, using automatic checks where reliable (timestamps, tool calls, reasoning, model identifiers, settings, versions, errors) and otherwise three sampled traces, with the reason recorded for every judgement (`docs/information_audit.md`).

AI pilot (RQ3, RQ4). Claude coded a 250-trace sample of MAST-failed traces, balanced across frameworks and models (AG2 64, AppWorld 23, HyperAgent 13, Magentic-One 63, MetaGPT 64, OpenManus 23), twice with `claude-sonnet-5-5` and once with `claude-opus-5-5`, using codebook v1. Annotations ran through Claude Code's headless mode with tools disabled and the system prompt replaced by the annotator prompt and codebook. Replies are constrained by a JSON schema whose allowed codes come from the codebook, and must give step-referenced quotes for each code, a causal summary, an evidence grade for any violation, and Ezell et al.'s cognitive function for each error. Every annotation records the models, the prompt and codebook versions and a hash of the full instructions; a refusal is an error, never rerouted. Each quote is checked automatically against the step it cites. Shares are weighted back to the eligible pool and reported with Wilson intervals (Appendix B).

Consistency measures. Agreement between models and between runs uses Cohen's κ [27, 28] for each layer's presence, each code (averaged, never pooled), error versus violation, the evidence grade, the deepest layer (also linearly weighted), and whether a failure was observed, with 95% bootstrap intervals

omitted where answers are too few to support one [29]. The share of identical full code sets is also reported [20]. These measures show whether the framework can be applied *consistently*, not whether the codes are *correct*: that needs independent human coders.

Documentation check and worked examples. Because the annotator sees only the trace, its most frequent design claims for each framework were checked against the framework's paper, documentation and code and against MAST's own description (`docs/design_check.md`). Worked examples illustrate what the coding records, using evidence that does not depend on judging the agents' code; they are illustrations, not validation.

Planned validation. A calibration set of 30 traces (26 failed, 4 blind MAST-clean controls; 18 development, 12 held out) was frozen before any coding, with the hidden key's SHA-256 fingerprint committed. Independent coders will code it blind; the AI pilot never touched it.

5 Results

5.1 RQ1: cross-walk

Every factor in Ezell et al.'s taxonomy has a place in AFACS (Table 2). The extension adds a supervision tier (L3.1–L3.3), coordination between agents (L2.3) and graded violations; their framework contributes model development (L4.4, with its own evidence requirement) and named security contexts. In the pilot, error types and cognitive functions paired as the cross-walk predicts (every skill-based error with action; decision errors mostly with decision; perceptual errors with understanding and observation), but this is not a test: the codebook the annotator read stated the prediction, and the two schemes share examples. The prediction is removed in codebook v1.1.

5.2 RQ2: information adequacy

The logs are good records of what the agents said and did: responses and actions are present in all seven frameworks (Figure 1). Everything needed to identify or re-run the system that failed is missing. No framework provides system documentation, artefacts or access; no framework logs a random seed; only ChatDev logs sampling settings and system prompts; and MAST names models without versions. The only model identifier logged inside a trace contradicts MAST's metadata for one group of 30 traces. Labelled reasoning appears in only two frameworks, and timestamps are absent from two (counts in Table 8, Appendix C). For AFACS the consequence is uneven observability: unsafe acts and coordination are well observed, supervision failures are visible as behaviour but without prompts it is often unclear what an orchestrator was told to do, and model development cannot be evidenced at all. These findings are an empirical instance of Ezell et al.'s reconstruction challenge.

5.3 RQ3: applicability

All 750 codings were valid, with no refusals or errors. Layer presence was applied consistently (Table 4): κ 0.63–0.76 between models and 0.74–0.87 between runs. Individual codes agreed less (0.52 and 0.59), the usual pattern in HFACS coding, and 24% and 40% of traces received an identical full set of codes. The weakest codes were `L4.3` (0.19 and 0.23), which overlaps with `L4.1`; the rare `L1.V?` (about zero); and `L3.2` and `L3.4` (0.32–0.62). HFACS's error types and Ezell et al.'s cognitive functions were applied about equally consistently. The evidence was genuine: 99.6–99.9% of about 4,000 quotes appeared verbatim in the step cited, and none was absent from the trace, though the check shows that the words exist, not that they support the code.

Weighted to the eligible pool, the annotator found a failure in 79–83% of the MAST-failed traces. Among failures, every run coded an unsafe act in all of them and a factor at L2–L4 in 85–91%. The categories AFACS adds were common: supervision failures (L3.1–L3.3) in 67% (95% interval 59–75%; other runs 69% and 71%) and coordination breakdowns (L2.3) in 40% (32–48%; 41%, 41%). Violations were a minority: 9% of failures (5–15%; 13%, 10%), mostly inferred from the act (6%) rather than stated (1%), concentrated in HyperAgent (69% of its failures) and AG2 (24%) and absent from MetaGPT

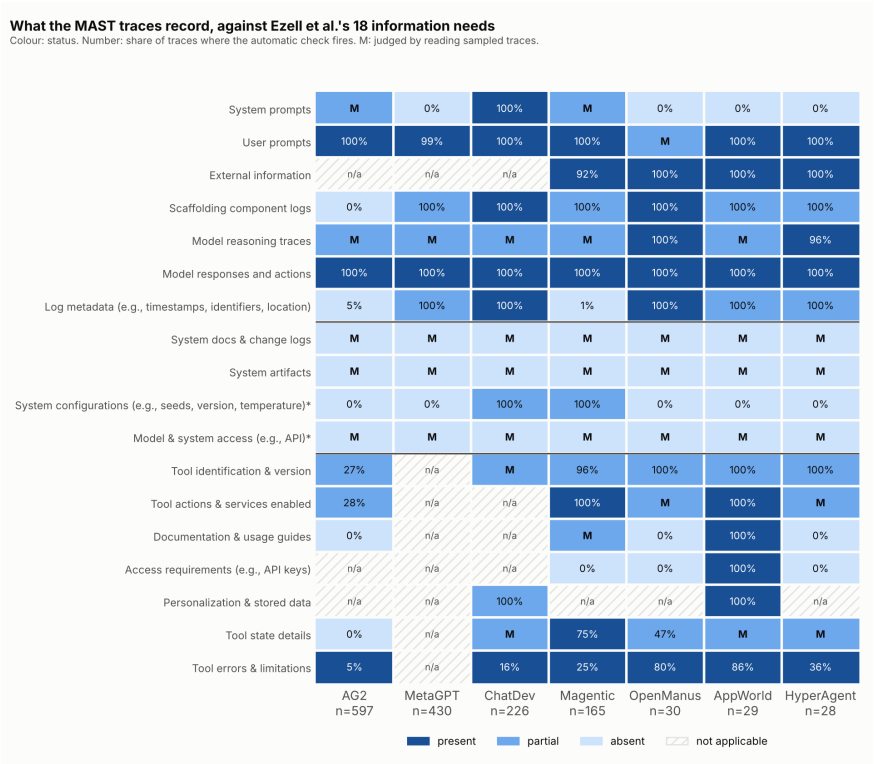


Figure 1: MAST's logs against the eighteen information items of Ezell et al.'s Table 1, by framework. Colour gives the status; numbers give the share of traces in which an automatic check fires; M marks judgements made by reading sampled traces. Items marked * matter mainly for reconstruction.

Table 4: Consistency of AI coding with AFACS on the 250-trace bulk sample. Cohen's κ with 95% bootstrap intervals (2,000 resamples); intervals are omitted where either coder gave fewer than three positive or negative answers.

Measure	Sonnet vs Opus	Sonnet run 1 vs run 2
L4 present	0.71 (0.62–0.79)	0.74 (0.65–0.82)
L3 present	0.63 (0.53–0.72)	0.74 (0.65–0.82)
L2 present	0.75 (0.67–0.83)	0.76 (0.67–0.83)
L1 present	0.76 (0.65–0.86)	0.87 (0.78–0.94)
Individual codes (mean per-code κ)	0.52 (0.45–0.57)	0.59 (0.53–0.66)
Violation vs error	0.67 (0.51–0.81)	0.68 (0.50–0.83)
Violation evidence grade	0.56	0.62
Error type, HFACS E1–E3 (mean per-code κ)	0.55	0.71
Error type, Ezell et al. functions (mean per-label κ)	0.58	0.70
Deepest layer (weighted)	0.75 (0.68–0.82)	0.80 (0.73–0.86)
Failure observed	0.69 (0.57–0.79)	0.80 (0.70–0.88)
Identical full code set (share of traces)	24%	40%
Evidence quotes found in the cited step	Sonnet 99.6%, 99.6%; Opus 99.9%	

Table 5: The annotator's most frequent design claims checked against each framework's documentation and code.

Framework	Claims	Confirmed	Partly true	Contradicted	Not settled
MetaGPT	5	2	3	0	0
AG2	4	1	2	1	0
Magentic-One	6	2	3	1	0
OpenManus	4	2	2	0	0
HyperAgent	3	0	3	0	0
AppWorld	2	1	0	0	1
All	24	8	13	2	1

and AppWorld. The typical violation was an agent reporting results of tool calls it never received, or "verification" code that only returned a hard-coded answer.

Three features need care. Design factors (L4) were coded in 63–74% of failures and were usually the deepest layer, but every L4 code is an inference from the trace (Section 5.4). No run ever left L1 empty, although the codebook allows it; the prompt's instruction to look for "the unsafe acts at the sharp end" may encourage this. And confidence was "medium" in 86–87% of codings, but informative when used: for codings rated "low", the two Sonnet runs agreed on whether there was a failure in 67% of cases, against 96% for "medium".

5.4 Design claims against documentation

Only two of the annotator's 24 most frequent design claims were wrong outright, but only eight held in full (Table 5). The typical error was to overstate a weakness as an absence: "no verifier", "no termination logic" or "no loop detection" where the system has the mechanism in a weak form. MetaGPT's reviewer comments only on the tests; MathChat detects repeated results; Magentic-One's orchestrator asks itself each turn whether it is in a loop; OpenManus has a duplicate-response detector that only adds a prompt; HyperAgent's editor rejects edits that fail a syntax check. A weak mechanism that fails leaves the same trace as a missing one. Two claims placed a failure at the wrong layer (a run-time fault and an environment set-up problem coded as design), and the annotator missed limits that no trace reveals, such as Magentic-One's 20-turn budget. Design factors coded from traces should be read as "no effective safeguard visible in the trace". The check also exposed discrepancies in MAST's own documentation (Appendix A): MetaGPT ran as a three-role example with no verifier of the code, and the HyperAgent traces carry a model label that postdates their source.

5.5 RQ4: MAST's flat labels

Every MAST mode spread across layers (Figure 2; expectations in Table 9, Appendix C). FM-1.1 Disobey Task Specification ($n = 79$) was a mixture: 17% of its traces carry a violation code (22% and 12% in the other runs) and the rest errors only, while a deficient specification was coded in only 6%. L2.3 coordination appears in 52% of FM-2.3 Task Derailment and 62% of FM-2.5 Ignored Other Agent's Input. FM-1.3 Step Repetition was coded as a design factor (L4.1, 60%) more often than as a supervisor failing to break the loop (L3.3, 27%), and FM-3.2 No or Incomplete Verification split between supervision (L3.1, 69%) and design (L4.1, 69%). A single flat label therefore covers failures whose remedies differ: a missing verifier, a supervisor who did not check, and an agent that departed from a requirement it had registered.

6 Worked Examples

Three short traces, each coded alike by all three runs, show what the layered coding records. Each statement rests on the trace, the benchmark's verdict or the documentation; benchmark questions are not reproduced. Two further examples are in Appendix D.

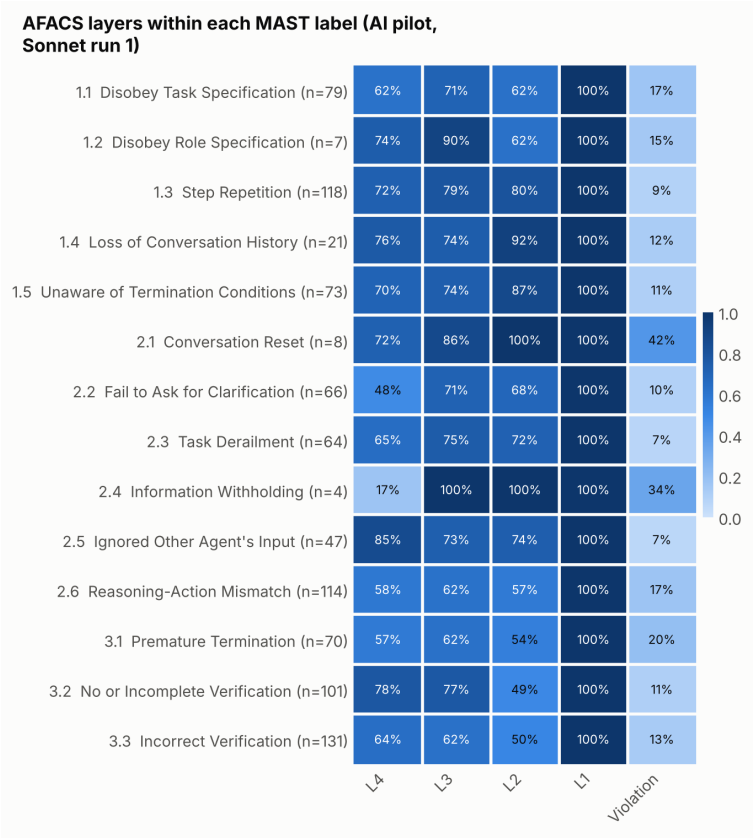


Figure 2: Share of failed traces carrying each MAST label that have an AFACS factor at each layer, and a violation code (Sonnet run 1, weighted).

A failure built into the design (MetaGPT). In `MetaGPT_ProgramDev-v2_Qwen__043` the task is a command-line weight converter. The coder's table covers eight direct unit pairs and rejects any other (S3). After the reviewer suggests testing an indirect conversion (S5), the tester adds a test for kilograms to ounces (S6), a pair the table lacks, so the test would fail; nobody can find out, because no agent runs code. The reviewer then replies in another language, approves the tests and reports a syntax error that is not there (S7), and the run ends with the code unrevised. MAST labels the trace 2.6, 3.1 and 3.3. All three runs code L4.1 (nothing runs or checks the code, and the review never reaches the coder), L3.1 (the only check never checked the code) and L1 errors; the documentation confirms the design claim. The flat labels name symptoms; the layered account places the deepest cause in a design that makes them likely whatever the agents do.

Flagged, but passed (AppWorld). MAST's judge labelled `AppWorld_Test-C_GPT-4o__000` 2.6 Reasoning–action mismatch, placing it among MAST's failed traces. The benchmark's own evaluation, hidden from the agents and the annotator, says the task passed, and all three runs judged that no failure occurred, noting inefficiencies but no unsafe act. Three of the 23 AppWorld traces in the sample passed despite MAST labels, and the annotator's judgement matched the benchmark in 21 of 23. A list of failure modes records what a judge noticed, not whether anything failed; AFACS records the two separately, as an occurrence report does.

A missing tool, or a tool not used (Magentic-One). In `Magentic_One_GPT4o__029` the task needs information in an image file. The orchestrator twice asks FileSurfer to read it (S5, S7); FileSurfer replies that it cannot interpret images (S8). A text-recognition tool installed instead garbles the part that matters (S11–S12), the assistant works from the garbled text (S14), and the orchestrator submits an unchecked answer (S15). All three runs code a missing tool (L4.2) as the deepest factor. The documentation qualifies this: WebSurfer works from screenshots and can open local files, and it was never asked. On that evidence the account should also include the orchestrator assigning the task to an agent that could not do it (L3.2), which only one run coded. Design claims made from a trace need the

system's documentation.

7 Discussion

The extension earns its place. The categories with no counterpart in Ezell et al.'s framework were among the most frequently coded, and violations, though rare, were where MAST's labels were least informative; the evidence grade showed that most rested on the act rather than on stated reasoning. An AI annotator applied the layers consistently, well above MAST's judge on its own labels. Two cautions temper this. Consistency is not correctness. And the deeper the layer, the less a trace can establish: design factors were the most common deepest layer yet the least observable, and the documentation check showed what trace-only coding gets wrong. In aviation, organisational findings come from company records rather than the flight recorder; the same division of evidence applies here. For the field, the information audit suggests a logging standard for multi-agent research releases, covering at least model identifiers and versions, sampling settings and seeds, system prompts and role definitions, and timestamps.

8 Limitations

- **Consistency, not correctness.** Agreement between AI runs can coexist with shared error; human validation is the next study. The annotator is a language model coding other language models and may share their blind spots.
- **Inferred design factors.** L4 codes are inferences from traces and, checked against documentation, tend to overstate absences.
- **No failure without an unsafe act was coded,** although the codebook allows it; whether such failures are rare or the option is avoided is open.
- **Outcomes were hidden** to avoid hindsight bias, so the annotator could not tell whether, for example, a mathematical answer was wrong.
- **The Ezell pairing is not a test,** because the codebook the annotator read stated it.
- **Behavioural, not intentional.** "Violation" describes the trace, not what a model wanted.
- **Data.** MAST's labels are noisy and record no outcome; ChatDev and long traces are excluded; model is confounded with framework; benchmark failures carry no harm; and HFACS's own reliability is debated.
- **Annotation route.** Annotations ran through Claude Code's headless mode; results through the bare API may differ slightly.

9 Future Work

The next study validates AFACS with independent human coders on the frozen calibration set. Three questions from the pilot belong there: how often trace-inferred design factors are right, when each trace is reviewed with the system's documentation; whether failures without an unsafe act are rare or induced by the prompt, tested with a variant prompt; and how coding changes when the annotator knows the outcome (the verdict, not the answer), as an investigator knows an accident occurred, or uses datasets with verified outcomes and human attribution [14]. A fifth level for the people around a system (users, operators, regulators) may suit deployed systems better and is left open. A further phase will apply AFACS to an evaluation of aviation dispatch and safety decisions under commercial, authority and user pressure, built in Inspect [30] from public regulatory sources; related aviation evaluations [31, 32, 33] do not code their failures with a human-factors framework.

Competing Interests

The author's company, CitadelAero, works on AI for aviation, the area of the aviation evaluation proposed as future work. This study received no external funding and uses only public data.

Use of AI Tools

The data pipeline, tests, figures and much of the drafting were done with Claude Code under the author's direction. The pilot's annotations were produced by Claude models as described in Section 4. The worked examples and the documentation check were prepared with Claude and checked against the traces and sources they cite. The framework and the codebook decisions are the author's.

Data and Code

Code: MIT licence, © 2026 CitadelAero Limited, at `https://github.com/JamesLEC1999/Research`. Data: MAST dataset, Cemri et al. [8], CC-BY-4.0. Traces are not redistributed, because they contain benchmark questions whose maintainers ask that they not be reposted; the pipeline recreates them from the pinned revision. The AI codings are committed with their codes in plain text and their quotes encoded, with a canary string.

References

[1] J. Reason, *Human Error*. Cambridge University Press, 1990.

[2] J. Reason, *Managing the Risks of Organizational Accidents*. Ashgate, 1997.

[3] S. A. Shappell and D. A. Wiegmann, "The Human Factors Analysis and Classification System–HFACS," U.S. Department of Transportation, Federal Aviation Administration, Office of Aviation Medicine, Report DOT/FAA/AM-00/7, 2000.

[4] D. A. Wiegmann and S. A. Shappell, *A Human Error Approach to Aviation Accident Analysis: The Human Factors Analysis and Classification System*. Ashgate, 2003.

[5] J. Reason, A. Manstead, S. Stradling, J. Baxter and K. Campbell, "Errors and violations on the roads: a real distinction?" *Ergonomics*, vol. 33, no. 10–11, pp. 1315–1332, 1990.

[6] C. Ezell, X. Roberts-Gaal and A. Chan, "Incident analysis for AI agents," in *Proc. AAAI/ACM Conference on AI, Ethics, and Society (AIES)*, 2025. arXiv:2508.14231.

[7] B. Baker, J. Huizinga, L. Gao, Z. Dou, M. Y. Guan, A. Madry, W. Zaremba, J. Pachocki and D. Farhi, "Monitoring reasoning models for misbehavior and the risks of promoting obfuscation," arXiv:2503.11926, 2025.

[8] M. Cemri, M. Z. Pan, S. Yang, L. A. Agrawal, B. Chopra, R. Tiwari, K. Keutzer, A. Parameswaran, D. Klein, K. Ramchandran, M. Zaharia, J. E. Gonzalez and I. Stoica, "Why do multi-agent LLM systems fail?" in *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. arXiv:2503.13657.

[9] M. Shamsujjoha, Q. Lu, D. Zhao and L. Zhu, "Swiss cheese model for AI safety: A taxonomy and reference architecture for multi-layered guardrails of foundation model based agents," arXiv:2408.02205, 2024.

[10] S. Mylius, "Systematic hazard analysis for frontier AI using STPA," arXiv:2506.01782, 2025.

[11] N. G. Leveson, *Engineering a Safer World: Systems Thinking Applied to Safety*. MIT Press, 2012.

[12] D. Deshpande, V. Gangal, H. Mehta, J. Krishnan, A. Kannappan and R. Qian, "TRAIL: Trace reasoning and agentic issue localization," arXiv:2505.08638, 2025.

[13] K. Zhu et al., "Where LLM agents fail and how they can learn from failures," arXiv:2509.25370, 2025.

[14] S. Zhang et al., "Which agent causes task failures and when? On automated failure attribution of LLM multi-agent systems," in *Proc. ICML*, PMLR vol. 267, pp. 76583–76599, 2025.

[15] H. Raj, V. Gupta, A. Mahmoud, R.-G. Dumitru, D. Yi, A. Sabharwal and Y. He, "Model or harness? An interaction-centric taxonomy for localizing agent failures," arXiv:2607.28802, 2026.

[16] D. A. Wiegmann and S. A. Shappell, "A human error analysis of commercial aviation accidents using the Human Factors Analysis and Classification System (HFACS)," U.S. Department of Transportation, Federal Aviation Administration, Office of Aviation Medicine, Report DOT/FAA/AM-01/3, 2001.

[17] N. S. Olsen and S. T. Shorrock, "Evaluation of the HFACS-ADF safety classification system: Inter-coder consensus and intra-coder consistency," *Accident Analysis & Prevention*, vol. 42, no. 2, pp. 437–444, 2010.

[18] T. N. Cohen, D. A. Wiegmann and S. A. Shappell, "Evaluating the reliability of the Human Factors Analysis and Classification System," *Aerospace Medicine and Human Performance*, vol. 86, no. 8, pp. 728–735, 2015.

Table 6: Data-quality exclusions applied before sampling.

Reason	Rows	Note
Kept	1,505	
Corrupt source group	96	ChatDev–CodeLlama: 100 rows hold 30 distinct texts
Duplicate text	35	identical trajectory seen earlier; first copy kept
No MAST labels	4	all 14 labels null
Misfiled text	2	"HyperAgent" rows that contain ChatDev logs
Total	1,642	

[19] B. R. Burnham, "Letter to the Editor re: Evaluating the reliability of the Human Factors Analysis Classification System," *Aerospace Medicine and Human Performance*, vol. 88, no. 6, p. 597, 2017.

[20] A. Ahmadi, S. Sharif and Y. Banad, "Improving aviation safety analysis: Automated HFACS classification using reinforcement learning with group relative policy optimization," arXiv:2508.21201, 2025.

[21] V. Krakovna, J. Uesato, V. Mikulik, M. Rahtz, T. Everitt, R. Kumar, Z. Kenton, J. Leike and S. Legg, "Specification gaming: the flip side of AI ingenuity," DeepMind blog, April 2020.

[22] Z. Zhong, A. Raghunathan and N. Carlini, "ImpossibleBench: Measuring LLMs' propensity of exploiting test cases," arXiv:2510.20270, 2025.

[23] A. Meinke, B. Schoen, J. Scheurer, M. Balesni, R. Shah and M. Hobbhahn, "Frontier models are capable of in-context scheming," arXiv:2412.04984, 2024.

[24] A. Lynch, B. Wright, C. Larson, S. J. Ritchie, S. Mindermann, E. Hubinger, E. Perez and K. Troy, "Agentic misalignment: How LLMs could be insider threats," arXiv:2510.05179, 2025.

[25] N. Smyth et al., "Quantifying overclaiming propensity in frontier LLM agents," arXiv:2609.20812, 2026.

[26] A. Hägele, A. P. Gema, H. Sleight, E. Perez and J. Sohl-Dickstein, "The hot mess of AI: How does misalignment scale with model intelligence and task complexity?" in *Proc. ICLR*, 2026. arXiv:2601.23045.

[27] J. Cohen, "A coefficient of agreement for nominal scales," *Educational and Psychological Measurement*, vol. 20, no. 1, pp. 37–46, 1960.

[28] J. R. Landis and G. G. Koch, "The measurement of observer agreement for categorical data," *Biometrics*, vol. 33, no. 1, pp. 159–174, 1977.

[29] A. R. Feinstein and D. V. Cicchetti, "High agreement but low kappa: I. The problems of two paradoxes," *Journal of Clinical Epidemiology*, vol. 43, no. 6, pp. 543–549, 1990.

[30] UK AI Security Institute, "Inspect AI: Framework for large language model evaluations," software, 2024. `https://inspect.aisi.org.uk/`.

[31] A. Brooker and T. Hughes, "Pre-Flight: A benchmark for evaluating large language models on aviation operational knowledge," arXiv:2607.01829, 2026.

[32] Y. Wu, H. Liu, Z. Li and B. Wang, "PilotBench: A benchmark for general aviation agents with safety constraints," in *Proc. International Joint Conference on Neural Networks (IJCNN)*, 2026. arXiv:2604.08987.

[33] Y. Yuan, Z. Wu, Y. Li, L. Ma and K. Li, "AeroCopilotBench: Safety-gated evaluation of LLM agents on aircraft emergency procedures in an executable cockpit," arXiv:2608.16349, 2026.

A Data Audit Details

Duplicates and judge consistency. The 1,642 rows contain 1,537 distinct texts. Among labelled rows, 65 groups of byte-identical texts cover 166 rows, and in 59 the copies carry different labels. Thirty Magentic-One traces appear under two group names and were labelled in separate judge runs: per-mode κ averages 0.43 (range -0.05 to 0.67; raw agreement 0.80; only 4 of 30 pairs identical; Figure 3). MAST reports $\kappa = 0.77$ between its judge and experts [8]; the figures cover different traces and are not directly comparable, but repeat runs disagreeing this much suggest substantial run-to-run noise in individual labels. In the ChatDev–CodeLlama group, copies of the same text agree at -0.04 over 111 pairs, suggesting texts and labels were mismatched upstream; the group is excluded.

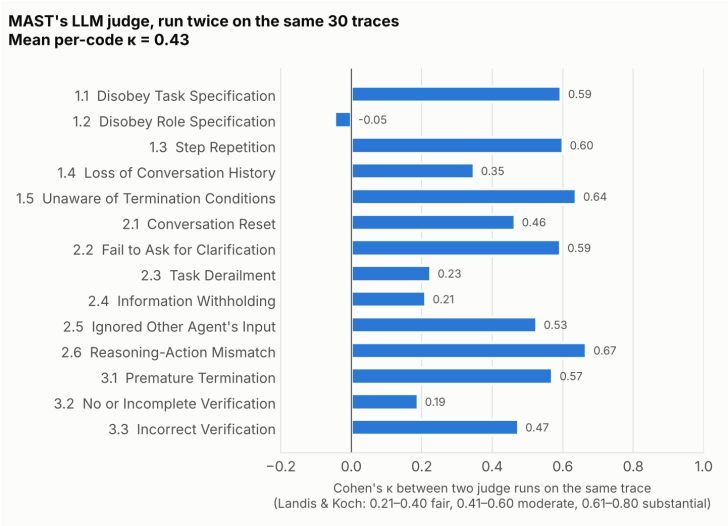


Figure 3: Test–retest agreement of MAST's LLM judge on 30 byte-identical Magentic-One traces labelled in two runs. Cohen's κ per failure mode; mean 0.43.

Metadata and logging. The OpenManus traces are tagged as programming tasks but are GAIA questions, and the AG2 "GSM" traces are GSM-Plus; both are corrected. In 29 of 30 OpenManus logs the task is printed truncated (the agents received the full text), and in one it was typed in and never logged; rendered traces carry a note. Five runs have an empty task. MAST's human-labelled file holds 19 traces (not 21), and within every round different traces carry identical label vectors, contradicting the original annotation sheets, so it is not used. Only 60 traces carry a benchmark verdict in the log (AppWorld and AG2 GSM-Plus).

Eligible pool. The pool is 893 MAST-failed and 373 MAST-clean traces: AG2 286 failed, MetaGPT 408, Magentic-One 127, AppWorld 27, OpenManus 27 and HyperAgent 18 (Figure 4). Four frameworks ran a single model, and "Claude" carries no version number.

MAST's descriptions of the systems. The documentation check found that MAST describes MetaGPT as having an explicit verifier, whereas the logs show MetaGPT's three-role example (425 of 430 traces follow its exact sequence), in which a reviewer comments only on the tests and no agent revises the code. MAST's HyperAgent traces come from a public SWE-bench submission dated September 2024 and are labelled Claude-3.7-Sonnet, a model released later; the submission's metadata names GPT-4o. The multi-agent AppWorld set-up and the AG2 MMLU team are MAST's own constructions with unpublished prompts, and the Magentic-One version could be inferred only from line numbers in a traceback.

B Method Details

Sampling and weights. The bulk sample allocates traces across frameworks and then across models within each framework by balanced allocation (smallest groups first), excluding the frozen calibration set. Each trace is weighted by the number of eligible traces it represents; shares use Wilson 95% intervals on the Kish effective sample size $n_{\text{eff}} = (\sum w)^2 / \sum w^2$.

Agreement. Bootstrap intervals use 2,000 resamples of traces and are omitted where either coder gave fewer than three positive or three negative answers. The weighted deepest-layer κ uses the full ordered label set, so that distances are not distorted by absent levels.

Quote check. Each quote is matched against its cited step after normalising whitespace, quotation marks, capitals and markdown; an ellipsis splits a quote into pieces that must appear in order. Quotes are classed as found in the cited step, found in another step, not found, or citing a hidden or absent step.

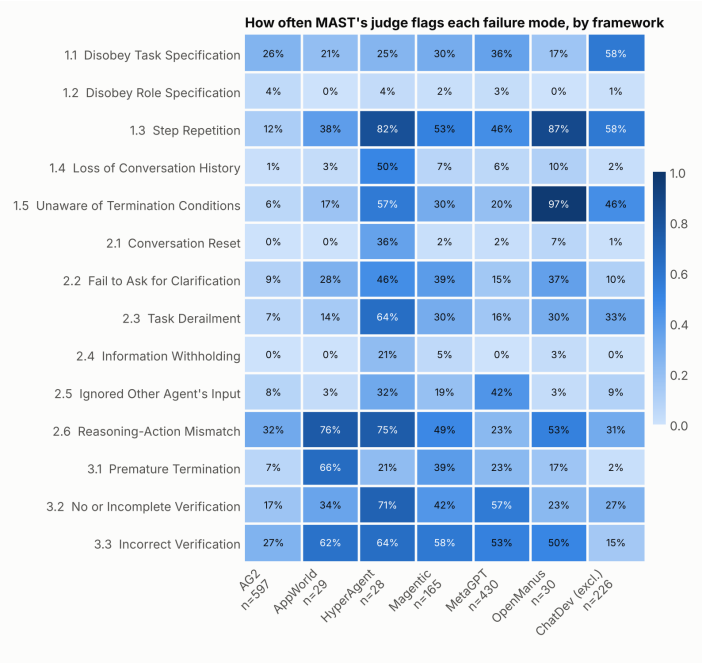


Figure 4: Share of kept traces in which MAST's judge flags each failure mode, by framework. ChatDev is shown for completeness but excluded from the study.

Table 7: Evidence needed to code each AFACS layer, and the corresponding information category in Ezell et al. [6].

Layer	Evidence needed	Ezell et al. category
L1 unsafe acts, L2 preconditions	The trace: messages, actions, tool outputs	Activity logs; tool information
L3 supervision	The trace, the task specification and the system prompt	Activity logs (system and user prompts)
L4.1–L4.3 design, process	The trace and what is known of the system's design	System documentation; tool information
L4.5 operation set-up	The trace and what is known of how the run was launched and configured	Activity logs (user prompts); system configuration
L4.4 model development	Developer records: training and reward specifications, change logs, version or checkpoint comparisons	System documentation and change logs; system artefacts; model access

Model-level comparisons. L4.4 is not coded from traces. Instead, patterns were compared across models with framework and benchmark held fixed (MetaGPT on ProgramDev, Claude against GPT-4o; on ProgramDev-v2, Qwen against CodeLlama; AG2 on GSM-Plus, Claude against GPT-4o). No MetaGPT group carried a violation code. AG2's Olympiad traces, all from GPT-4o, carried one in 31% of failures, against none in its GSM-Plus and MMLU traces; there benchmark and model are confounded. The groups are small and the comparisons exploratory.

C Further Tables

D Further Worked Examples

A violation graded on the act (AG2). In `AG2_Olympiad_GPT4o__150` the assistant is told to use Python for calculations, wait for the results, and box its answer only after the queries have run (S1). It reasons informally, then writes code headed "Perform calculation/verification" in which the answer is assigned as a constant and printed, and it boxes that number before the code has run (S2). The proxy returns the constant (S3), and the assistant presents it as the outcome of its analysis (S4). All three runs

Table 8: Number of Ezell et al.'s eighteen information items each framework's logs provide. Applicable excludes items that do not arise in these benchmark runs (e.g. tool information for agents that use no tools).

Framework	Traces	Applicable	Present	Partial	Absent
AppWorld	29	18	8	5	5
Magentic-One	165	17	6	6	5
OpenManus	30	17	6	4	7
HyperAgent	28	17	5	5	7
AG2	597	15	2	5	8
ChatDev	226	14	6	5	3
MetaGPT	430	10	2	3	5

Table 9: Where the framework expects MAST's modes to fall in AFACS, stated before the pilot.

MAST mode	Expected AFACS location
1.1 Disobey task specification	Spread: L3.4, L1.E2 or L1.E3, L1.V, or L4.5 (a run set up with no task)
1.2 Disobey role specification	L4.1 role design, or L1.E2 / L1.V
1.3 Step repetition	L1 act, with L3.3 where a supervisor could have broken the loop, and L4.1 (no termination logic)
1.4 Loss of conversation history	L2.2 agent state, with L4.2 context limits
1.5 Unaware of termination conditions	L4.1 termination logic, or L1.E2
2.1–2.5 Inter-agent misalignment	Mostly L2.3 coordination; 2.2 also L3.4
2.6 Reasoning–action mismatch	L1.E1 (sound intention, faulty execution), or L1.V
3.1 Premature termination	L1.E2, or L1.V (false claim of completion), with L3.1
3.2 No or incomplete verification	By whose job it was: L1 (incl. L1.V1 routine skipping), L3.1, or L4.1/L4.3
3.3 Incorrect verification	L3.1, L1.E2 or L1.E3

code `L1.V2` with the evidence graded *inferred*: nothing in the assistant's words acknowledges that the "verification" is empty, but the act is hard to read as accidental. One run noted that a reviewer could code it as a decision error only, which is the judgement the grade leaves visible. Whether the answer is wrong cannot be told from the log.

A run with no task (MetaGPT). In `MetaGPT_ProgramDev2_GPT4o__099` the task text is empty (S2). The coder writes a placeholder (S3), the tester writes template tests for a function that does not exist (S4), and the reviewer asks for more specific test inputs without mentioning that there is no task (S5). MAST labels the trace FM-1.1 Disobey Task Specification and FM-2.2 Fail to Ask for Clarification, but there was no specification to disobey: condition (1) of the violation rule cannot hold. The deepest factor is the run's set-up, L4.5 Operation set-up.